

YASH DHUMAL C3 BATCH 14369

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv("uber.csv")
df
```

```
Out[2]:
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude
0	24238194	52:06.0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	
1	27835199	04:56.0	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	
2	44984355	45:00.0	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	
3	25894730	22:21.0	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	
4	17610152	47:00.0	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	
...
199995	42598914	49:00.0	3.0	2012-10-28 10:49:00 UTC	-73.987042	40.739367	
199996	16382965	09:00.0	7.5	2014-03-14 01:09:00 UTC	-73.984722	40.736837	
199997	27804658	42:00.0	30.9	2009-06-29 00:42:00 UTC	-73.986017	40.756487	
199998	20259894	56:25.0	14.5	2015-05-20 14:56:25 UTC	-73.997124	40.725452	
199999	11951496	08:00.0	14.1	2010-05-15 04:08:00 UTC	-73.984395	40.720077	

200000 rows × 9 columns

```
In [3]: df.head()
```

Out[3]:

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
--	------------	-----	-------------	-----------------	------------------	-----------------	-------------------	------------------	-----------------

0	24238194	52:06.0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999817	40.738354	1
1	27835199	04:56.0	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994355	40.728225	1
2	44984355	45:00.0	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-74.005043	40.740770	1
3	25894730	22:21.0	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.976124	40.790844	1
4	17610152	47:00.0	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.925023	40.744085	1

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Unnamed: 0            200000 non-null int64  
1   key                   200000 non-null object 
2   fare_amount           200000 non-null float64 
3   pickup_datetime       200000 non-null object 
4   pickup_longitude      200000 non-null float64 
5   pickup_latitude       200000 non-null float64 
6   dropoff_longitude     199999 non-null float64 
7   dropoff_latitude     199999 non-null float64 
8   passenger_count       200000 non-null int64  
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

In [5]: `df.columns`

Out[5]: Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime', 'pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude', 'passenger_count'], dtype='object')

In [6]: `df = df.drop(['Unnamed: 0', 'key'], axis = 1)`

In [7]: `df.head()`

Out[7]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	40.7
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	40.7
2	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	40.7
3	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	40.8
4	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	40.7

In [8]: `df.shape`

Out[8]: (200000, 7)

In [9]: `df.dtypes`

Out[9]:

fare_amount	float64
pickup_datetime	object
pickup_longitude	float64
pickup_latitude	float64
dropoff_longitude	float64
dropoff_latitude	float64
passenger_count	int64
dtype:	object

In [10]: `df.describe()`

Out[10]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
count	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000	200000.000000
mean	11.359955	-72.527638	39.935885	-72.525292	39.923890	1.716961
std	9.901776	11.437787	7.720539	13.117408	6.794829	1.227004
min	-52.000000	-1340.648410	-74.015515	-3356.666300	-881.985513	0.000000
25%	6.000000	-73.992065	40.734796	-73.991407	40.733823	1.000000
50%	8.500000	-73.981823	40.752592	-73.980093	40.753042	1.000000
75%	12.500000	-73.967153	40.767158	-73.963659	40.768001	1.000000
max	499.000000	57.418457	1644.421482	1153.572603	872.697628	16.000000

In [11]: `df.isnull().sum()`

Out[11]:

fare_amount	0
pickup_datetime	0
pickup_longitude	0
pickup_latitude	0
dropoff_longitude	1
dropoff_latitude	1
passenger_count	0
dtype:	int64

```
In [12]: df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(), inplace=True)
```

```
In [13]: df.isnull().sum()
```

```
Out[13]: fare_amount      0
pickup_datetime      0
pickup_longitude     0
pickup_latitude      0
dropoff_longitude     1
dropoff_latitude      0
passenger_count      0
dtype: int64
```

```
In [14]: df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].median(), inplace=True)
```

```
In [15]: df.isnull().sum()
```

```
Out[15]: fare_amount      0
pickup_datetime      0
pickup_longitude     0
pickup_latitude      0
dropoff_longitude     0
dropoff_latitude      0
passenger_count      0
dtype: int64
```

```
In [16]: df.dtypes
```

```
Out[16]: fare_amount      float64
pickup_datetime      object
pickup_longitude     float64
pickup_latitude      float64
dropoff_longitude     float64
dropoff_latitude      float64
passenger_count      int64
dtype: object
```

```
In [17]: df.pickup_datetime = pd.to_datetime(df.pickup_datetime, errors='coerce')
```

```
In [18]: df.dtypes
```

```
Out[18]: fare_amount      float64
pickup_datetime      datetime64[ns, UTC]
pickup_longitude     float64
pickup_latitude      float64
dropoff_longitude     float64
dropoff_latitude      float64
passenger_count      int64
dtype: object
```

```
In [19]: df = df.assign(hour = df.pickup_datetime.dt.hour,
                        day = df.pickup_datetime.dt.day,
                        month = df.pickup_datetime.dt.month,
                        year = df.pickup_datetime.dt.year,
                        dayofweek = df.pickup_datetime.dt.dayofweek)
```

```
In [20]: df.head
```

```
Out[20]: <bound method NDFrame.head of
p_longitude \
0          7.5 2015-05-07 19:52:06+00:00      -73.999817
1          7.7 2009-07-17 20:04:56+00:00      -73.994355
2         12.9 2009-08-24 21:45:00+00:00      -74.005043
3          5.3 2009-06-26 08:22:21+00:00      -73.976124
4         16.0 2014-08-28 17:47:00+00:00      -73.925023
...
199995      3.0 2012-10-28 10:49:00+00:00      -73.987042
199996      7.5 2014-03-14 01:09:00+00:00      -73.984722
199997     30.9 2009-06-29 00:42:00+00:00      -73.986017
199998     14.5 2015-05-20 14:56:25+00:00      -73.997124
199999     14.1 2010-05-15 04:08:00+00:00      -73.984395
```

```

pickup_latitude dropoff_longitude dropoff_latitude passenger_count \
0          40.738354      -73.999512          40.723217          1
1          40.728225      -73.994710          40.750325          1
2          40.740770      -73.962565          40.772647          1
3          40.790844      -73.965316          40.803349          3
4          40.744085      -73.973082          40.761247          5
...
199995      40.739367      -73.986525          40.740297          1
199996      40.736837      -74.006672          40.739620          1
199997      40.756487      -73.858957          40.692588          2
199998      40.725452      -73.983215          40.695416          1
199999      40.720077      -73.985508          40.768793          1
```

```

hour  day  month  year  dayofweek
0      19    7     5  2015         3
1      20   17     7  2009         4
2      21   24     8  2009         0
3       8   26     6  2009         4
4      17   28     8  2014         3
...
199995   10   28    10  2012         6
199996    1   14     3  2014         4
199997    0   29     6  2009         0
199998   14   20     5  2015         2
199999    4   15     5  2010         5
```

[200000 rows x 12 columns]>

```
In [21]: df = df.drop('pickup_datetime', axis =1)
```

```
In [22]: df.head
```

```
Out[22]: <bound method NDFrame.head of
de dropoff_longitude \
0          7.5          -73.999817          40.738354          -73.999512
1          7.7          -73.994355          40.728225          -73.994710
2         12.9          -74.005043          40.740770          -73.962565
3          5.3          -73.976124          40.790844          -73.965316
4         16.0          -73.925023          40.744085          -73.973082
...         ...         ...         ...         ...
199995      3.0          -73.987042          40.739367          -73.986525
199996      7.5          -73.984722          40.736837          -74.006672
199997     30.9          -73.986017          40.756487          -73.858957
199998     14.5          -73.997124          40.725452          -73.983215
199999     14.1          -73.984395          40.720077          -73.985508

dropoff_latitude passenger_count hour day month year dayofweek
0          40.723217             1   19    7     5  2015         3
1          40.750325             1   20   17     7  2009         4
2          40.772647             1   21   24     8  2009         0
3          40.803349             3    8   26     6  2009         4
4          40.761247             5   17   28     8  2014         3
...         ...         ...   ...   ...   ...   ...         ...
199995      40.740297             1   10   28    10  2012         6
199996      40.739620             1    1   14     3  2014         4
199997      40.692588             2    0   29     6  2009         0
199998      40.695416             1   14   20     5  2015         2
199999      40.768793             1    4   15     5  2010         5

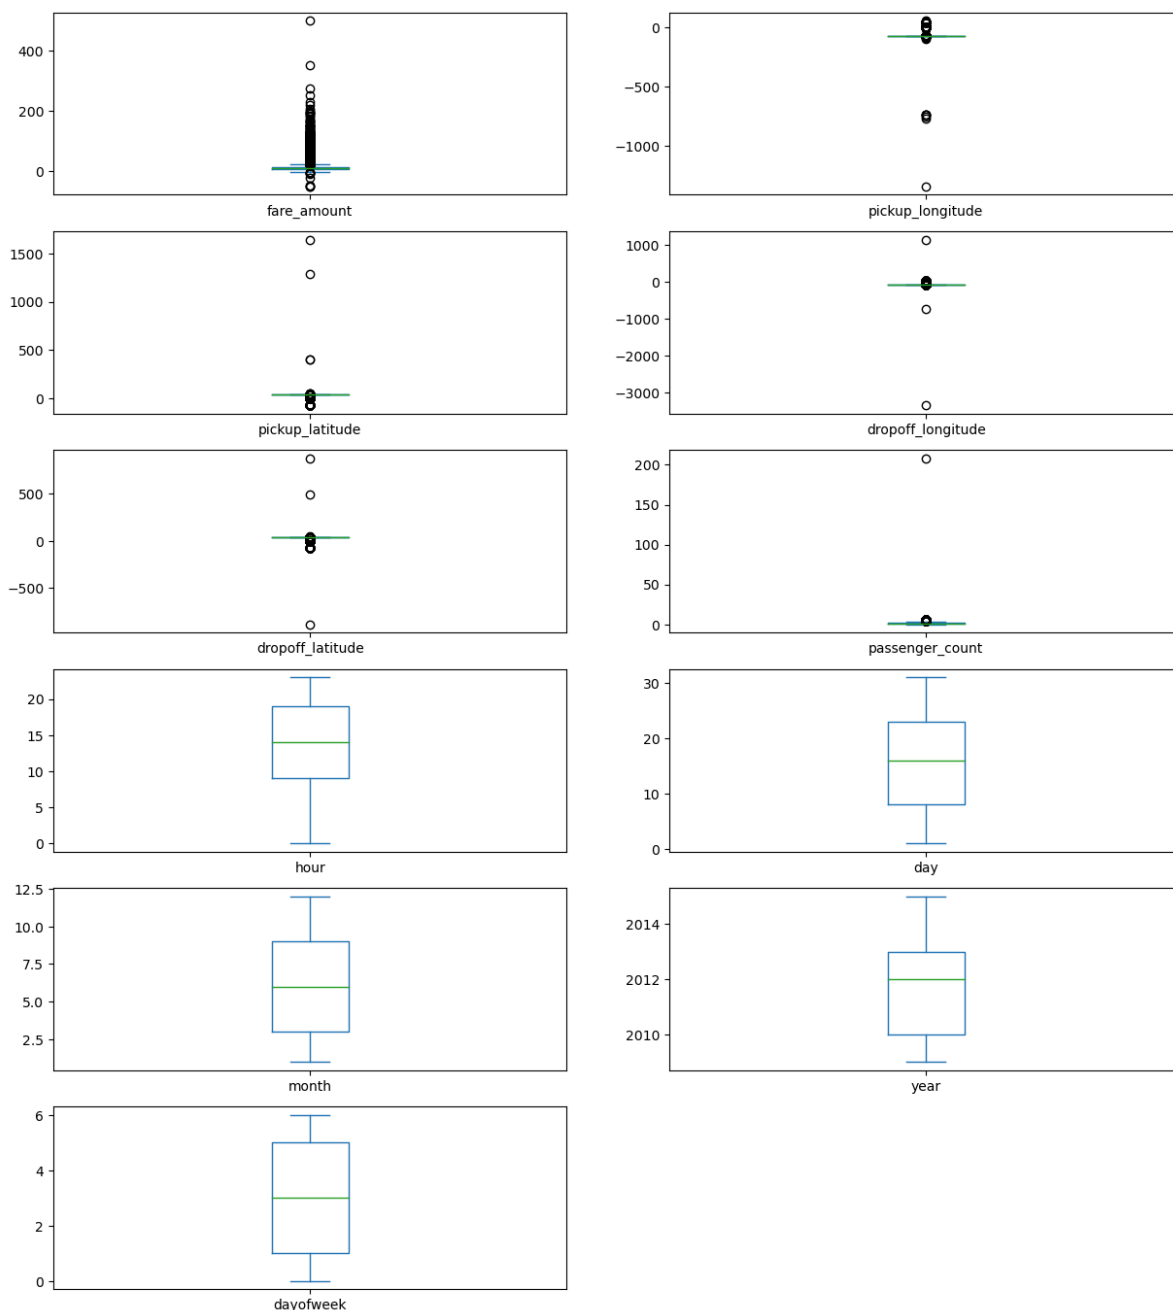
[200000 rows x 11 columns]>
```

```
In [23]: df.dtypes
```

```
Out[23]: fare_amount          float64
pickup_longitude         float64
pickup_latitude          float64
dropoff_longitude         float64
dropoff_latitude          float64
passenger_count           int64
hour                      int32
day                       int32
month                     int32
year                      int32
dayofweek                 int32
dtype: object
```

```
In [25]: df.plot(kind="box", subplots = True, layout =(7,2),figsize=(15,20))
```

```
Out[25]: fare_amount          Axes(0.125,0.786098;0.352273x0.0939024)
pickup_longitude         Axes(0.547727,0.786098;0.352273x0.0939024)
pickup_latitude          Axes(0.125,0.673415;0.352273x0.0939024)
dropoff_longitude         Axes(0.547727,0.673415;0.352273x0.0939024)
dropoff_latitude          Axes(0.125,0.560732;0.352273x0.0939024)
passenger_count           Axes(0.547727,0.560732;0.352273x0.0939024)
hour                      Axes(0.125,0.448049;0.352273x0.0939024)
day                       Axes(0.547727,0.448049;0.352273x0.0939024)
month                     Axes(0.125,0.335366;0.352273x0.0939024)
year                      Axes(0.547727,0.335366;0.352273x0.0939024)
dayofweek                 Axes(0.125,0.222683;0.352273x0.0939024)
dtype: object
```



```
In [27]: def remove_outlier(df1, col):
          Q1 = df1[col].quantile(0.25)
          Q3 = df1[col].quantile(0.75)
          IQR = Q3 - Q1
          lower_whisker = Q1 - 1.5 * IQR
          upper_whisker = Q3 + 1.5 * IQR
          df[col] = np.clip(df1[col], lower_whisker, upper_whisker)
          return df1
          def treat_outliers_all(df1, col_list):
              for c in col_list:
                  df1 = remove_outlier(df1, c)
              return df1
```

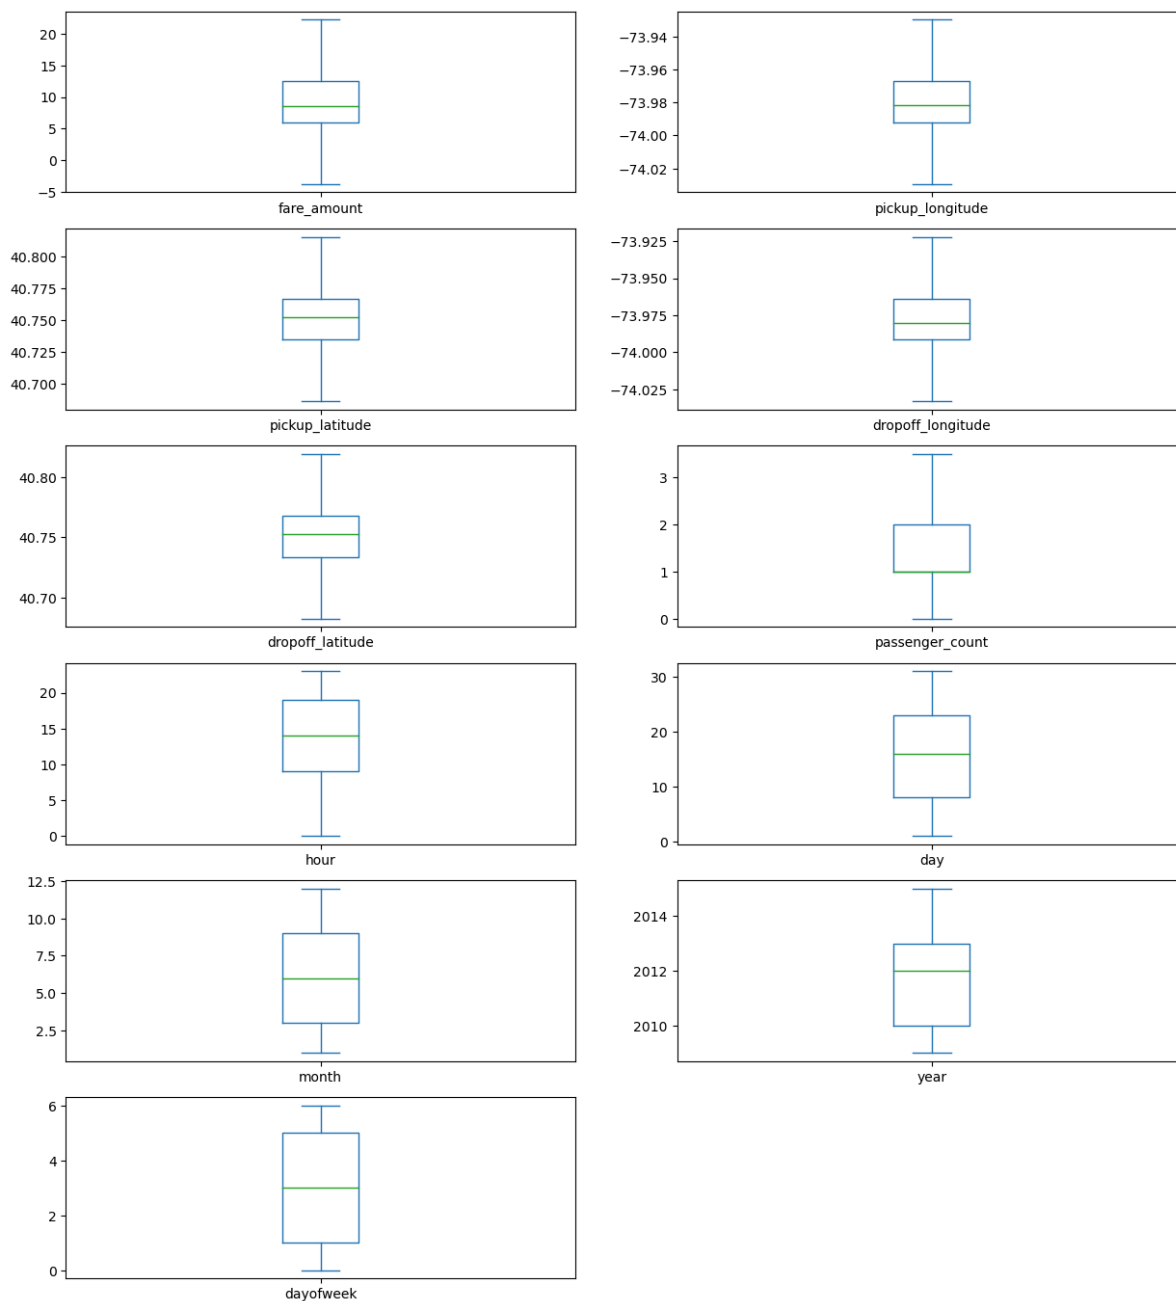
```
In [28]: df = treat_outliers_all(df, df.iloc[:, 0::])
```

```
In [29]: df.plot(kind = "box", subplots = True, layout=(7,2), figsize=(15,20))
```

```

Out[29]: fare_amount      Axes(0.125,0.786098;0.352273x0.0939024)
pickup_longitude Axes(0.547727,0.786098;0.352273x0.0939024)
pickup_latitude  Axes(0.125,0.673415;0.352273x0.0939024)
dropoff_longitude Axes(0.547727,0.673415;0.352273x0.0939024)
dropoff_latitude  Axes(0.125,0.560732;0.352273x0.0939024)
passenger_count  Axes(0.547727,0.560732;0.352273x0.0939024)
hour             Axes(0.125,0.448049;0.352273x0.0939024)
day             Axes(0.547727,0.448049;0.352273x0.0939024)
month           Axes(0.125,0.335366;0.352273x0.0939024)
year           Axes(0.547727,0.335366;0.352273x0.0939024)
dayofweek       Axes(0.125,0.222683;0.352273x0.0939024)
dtype: object

```



```

In [30]: pip install haversine

```


Collecting haversine

Obtaining dependency information for haversine from <https://files.pythonhosted.org/packages/5b/f1/b7274966f0b5b665d9114e86d09c6bc87d241781d63d8817323dcfa940c6/haversine-2.8.1-py2.py3-none-any.whl.metadata>

Downloading haversine-2.8.1-py2.py3-none-any.whl.metadata (5.9 kB)

Downloading haversine-2.8.1-py2.py3-none-any.whl (7.7 kB)

Installing collected packages: haversine

Successfully installed haversine-2.8.1

Note: you may need to restart the kernel to use updated packages.

```
In [38]: import haversine as hs
travel_dist = []
for pos in range(len(df['pickup_longitude'])):
    long1,lati1,long2,lati2 = [df['pickup_longitude'][pos],df['pickup_latitude'][pos],df['dropoff_longitude'][pos],df['dropoff_latitude'][pos]]
    loc1 =(lati1,long1)
    loc2 =(lati2,long2)
    c = hs.haversine(loc1,loc2)
    travel_dist.append(c)

print(travel_dist)
df['dist_travel_km']= travel_dist
df.head
```

IOPub data rate exceeded.

The notebook server will temporarily stop sending output to the client in order to avoid crashing it.

To change this limit, set the config variable

`--NotebookApp.iopub_data_rate_limit`.

Current values:

NotebookApp.iopub_data_rate_limit=1000000.0 (bytes/sec)

NotebookApp.rate_limit_window=3.0 (secs)

```
Out[38]: <bound method NDFrame.head of
de dropoff_longitude \
0          7.50      -73.999817      40.738354      -73.999512
1          7.70      -73.994355      40.728225      -73.994710
2         12.90      -74.005043      40.740770      -73.962565
3          5.30      -73.976124      40.790844      -73.965316
4         16.00      -73.929786      40.744085      -73.973082
...         ...         ...         ...         ...
199995      3.00      -73.987042      40.739367      -73.986525
199996      7.50      -73.984722      40.736837      -74.006672
199997     22.25      -73.986017      40.756487      -73.922036
199998     14.50      -73.997124      40.725452      -73.983215
199999     14.10      -73.984395      40.720077      -73.985508

dropoff_latitude passenger_count hour day month year dayofweek \
0          40.723217           1.0   19   7     5  2015          3
1          40.750325           1.0  20  17     7  2009          4
2          40.772647           1.0  21  24     8  2009          0
3          40.803349           3.0   8  26     6  2009          4
4          40.761247           3.5  17  28     8  2014          3
...         ...         ...   ...   ...   ...   ...         ...
199995      40.740297           1.0  10  28    10  2012          6
199996      40.739620           1.0   1  14     3  2014          4
199997      40.692588           2.0   0  29     6  2009          0
199998      40.695416           1.0  14  20     5  2015          2
199999      40.768793           1.0   4  15     5  2010          5

dist_travel_km
0          1.683325
1          2.457593
2          5.036384
3          1.661686
4          4.116088
...         ...
199995      0.112210
199996      1.875053
199997      8.919323
199998      3.539720
199999      5.417791
```

[200000 rows x 12 columns]>

```
In [39]: df = df.loc[(df.dist_travel_km>=1) | (df.dist_travel_km<=130) ]
print("Remaining obervation:" , df.shape)
```

Remaining obervation: (200000, 12)

```
In [47]: incorrect_coordinates =df.loc[(df.pickup_latitude>90) | (df.pickup_latitude< -90)|
(df.dropoff_latitude>90) | (df.dropoff_latitude< -90)|
(df.pickup_longitude>180) | (df.pickup_longitude< -180)|
(df.dropoff_longitude>90) | (df.dropoff_longitude< -90)
]
```

```
In [48]: df.drop(incorrect_coordinates, inplace = True, errors ='ignore')
```

```
In [49]: df.head()
```

Out[49]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_
0	7.5	-73.999817	40.738354	-73.999512	40.723217	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	
4	16.0	-73.929786	40.744085	-73.973082	40.761247	

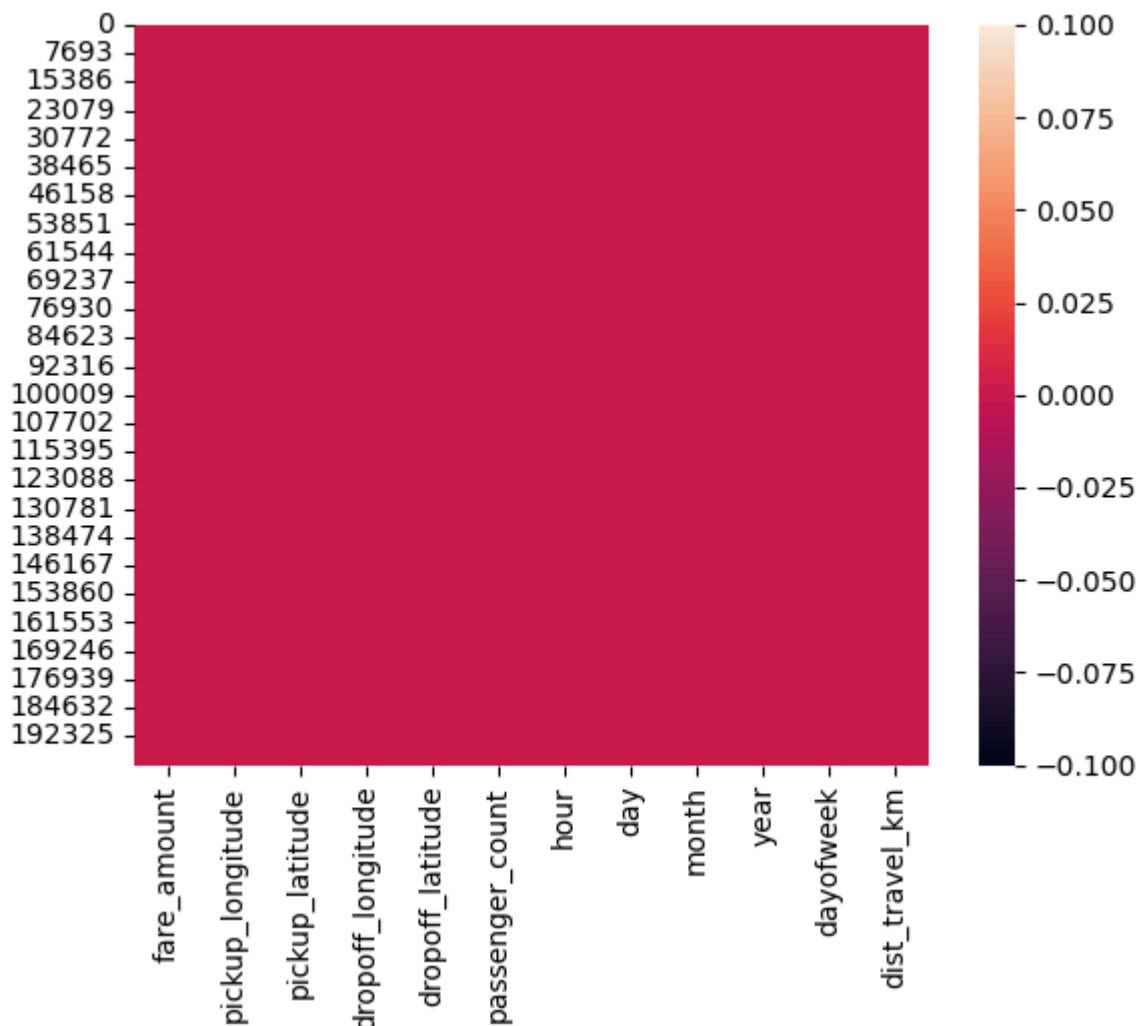
In [50]: `df.isnull().sum()`

Out[50]:

fare_amount	0
pickup_longitude	0
pickup_latitude	0
dropoff_longitude	0
dropoff_latitude	0
passenger_count	0
hour	0
day	0
month	0
year	0
dayofweek	0
dist_travel_km	0
dtype: int64	

In [51]: `sns.heatmap(df.isnull())`

Out[51]: <Axes: >



```
In [52]: corr = df.corr()
```

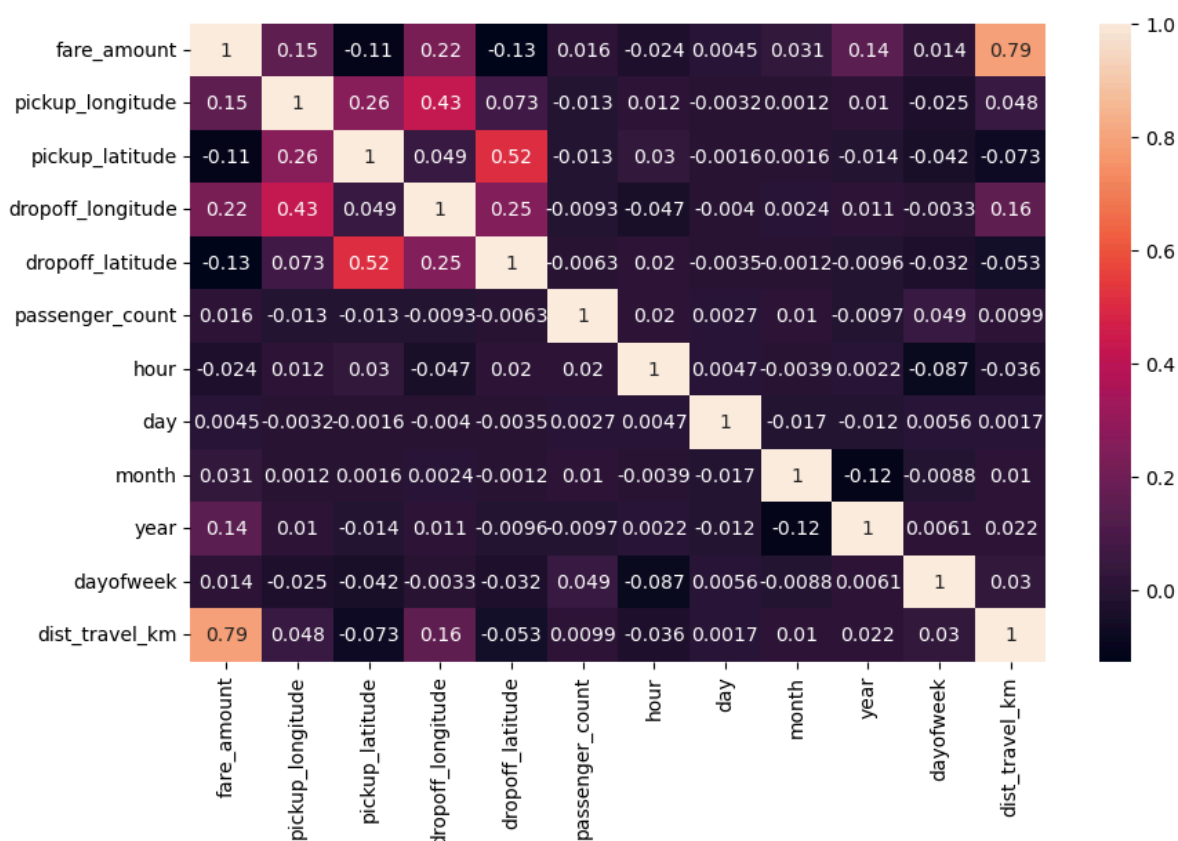
```
In [53]: corr
```

```
Out[53]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
fare_amount	1.000000	0.154069	-0.110842	0.218675	-0.125898
pickup_longitude	0.154069	1.000000	0.259497	0.425619	0.073290
pickup_latitude	-0.110842	0.259497	1.000000	0.048889	0.515714
dropoff_longitude	0.218675	0.425619	0.048889	1.000000	0.245667
dropoff_latitude	-0.125898	0.073290	0.515714	0.245667	1.000000
passenger_count	0.015778	-0.013213	-0.012889	-0.009303	-0.006144
hour	-0.023623	0.011579	0.029681	-0.046558	0.019100
day	0.004534	-0.003204	-0.001553	-0.004007	-0.003204
month	0.030817	0.001169	0.001562	0.002391	-0.001169
year	0.141277	0.010198	-0.014243	0.011346	-0.009303
dayofweek	0.013652	-0.024652	-0.042310	-0.003336	-0.031000
dist_travel_km	0.786385	0.048446	-0.073362	0.155191	-0.052000

```
In [55]: fig,axis = plt.subplots(figsize= (10,6))
sns.heatmap(df.corr(), annot = True)
```

```
Out[55]: <Axes: >
```



```
In [56]: df.dtypes
```

```
Out[56]: fare_amount          float64
pickup_longitude      float64
pickup_latitude       float64
dropoff_longitude     float64
dropoff_latitude      float64
passenger_count       float64
hour                  int32
day                   int32
month                 int32
year                  int32
dayofweek             int32
dist_travel_km        float64
dtype: object
```

```
In [58]: x = df[['pickup_longitude', 'pickup_latitude','dropoff_longitude','dropoff_latitude']]
```

```
In [59]: y = df['fare_amount']
```

```
In [60]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train,y_test = train_test_split(x,y,test_size = 0.33)
```

```
In [61]: from sklearn.linear_model import LinearRegression
regression = LinearRegression()
```

```
In [62]: regression.fit(x_train,y_train)
```

```
Out[62]: ▼ LinearRegression
LinearRegression()
```

```
In [64]: regression.intercept_
```

```
Out[64]: 3615.5861559214154
```

```
In [65]: regression.coef_
```

```
Out[65]: array([ 2.49068742e+01, -6.89454719e+00,  2.01504052e+01, -1.81286709e+01,
        5.98985068e-02,  6.08043699e-03,  2.92131129e-03,  5.97438413e-02,
        3.68715386e-01, -3.72169236e-02,  1.85329279e+00])
```

```
In [66]: prediction = regression.predict(x_test)
```

```
In [67]: print(prediction)
```

```
[ 9.21179204  8.91436439  5.47686097 ...  7.52253896 14.2121657
 5.73971605]
```

```
In [68]: y_test
```

```
Out[68]: 134041    13.5
125453     8.5
151639     3.0
174239    12.9
165574     5.7
...
73077     4.9
91779    14.0
176467     5.7
152546    17.7
194320     6.1
Name: fare_amount, Length: 66000, dtype: float64
```

```
In [69]: from sklearn.metrics import r2_score
```

```
In [70]: r2_score(y_test, prediction)
```

```
Out[70]: 0.6599227836474304
```

```
In [71]: from sklearn.metrics import mean_squared_error
```

```
In [72]: MSE = mean_squared_error(y_test, prediction)
```

```
In [73]: MSE
```

```
Out[73]: 10.009319812626584
```